

Sean McLean

ALY 6015

Module 4 Assignment

Regularization

Introduction


The mission of this assignment is to show how a dependent variable can be predicted with the usage of regularization models. We will be using the same dataset as the previous assignment to demonstrate which of the used models like Ridge and Lasso will perform better. The previous dataset used is called 'college' which contains 18 variables and 777 rows, and the predicted variable we will use is called 'Grad.rate'. After looking at the performances of both the lasso and ridge models, a stepwise model will be created to see if it has better results than the previous modeling techniques.

Analysis: Ridge Regression

1. The first step is to break up the college data set into training and test sets as was previously done in the last assignment. This is a customary practice in machine learning where the dataset is divided into a pair of subsets with the training data going through a process where the model identifies and learns patterns and trends in the data. The training data like the previous assignment will also be 70 percent of the dataset. The remaining portion will be the test set that will be used to look at the model's performance on data not encountered before (Barkved, 2022). The two different subsets are put into variables and then matrix variables where the 'Grad.rate' variable is removed and put into a separate variable for regularization modeling.

2. and 3. The ridge regression model will be built first by using the `cv.glmnet` function to establish the `lambda.min` and `lambda.lse` values and then will plot them. The `lambda` parameter is important in the regularization process, performing tasks like keeping overfitting under control and shrinking coefficients.

```
> set.seed(123)
> cv.ridge <- cv.glmnet(train_x, train_y, alpha = 0, nfolds = 10)
> log(cv.ridge$lambda.min)
[1] 0.3880099
> log(cv.ridge$lambda.1se)
[1] 2.62082
```



The lambda.min value is 0.388 which indicates that it could be less sparse with regards to the coefficients in the model. The lambda.1se value of 2.621 means that it could be sparser and the coefficients being smaller. In the plot the vertical dotted lines show where the lambda.min and lambda.1se values fall, and the number of variables from the data set all remain which could potentially mean that they are all still relevant to the model.

4. Using the cv.glmnet function, a ridge regression model is fitted, and the coefficients are compared between lambda.min and lambda.1se followed by the coefficients with no regularization.

<code>> coef(model_1se_ridge)</code>		<code>> coef(model_1se_ridge)</code>	
18 x 1 sparse Matrix of class "dgCMatrix"		18 x 1 sparse Matrix of class "dgCMatrix"	
s0		s0	
(Intercept)	3.972866e+01	(Intercept)	4.189904e+01
PrivateYes	4.373132e+00	PrivateYes	3.293228e+00
Apps	8.612021e-04	Apps	2.940924e-04
Accept	1.632762e-04	Accept	2.340064e-04
Enroll	-3.615392e-04	Enroll	3.715852e-05
Top10perc	6.480415e-02	Top10perc	8.291118e-02
Top25perc	1.426209e-01	Top25perc	9.466930e-02
F.Undergrad	-4.552089e-05	F.Undergrad	-4.208574e-05
P.Undergrad	-1.185306e-03	P.Undergrad	-8.630806e-04
Outstate	6.398481e-04	Outstate	4.870042e-04
Room.Board	2.043375e-03	Room.Board	1.492526e-03
Books	-5.169888e-04	Books	-6.175465e-04
Personal	-2.939102e-03	Personal	-2.300453e-03
PhD	1.004773e-01	PhD	4.525940e-02
Terminal	-9.560248e-02	Terminal	4.658045e-04
S.F.Ratio	-1.801664e-01	S.F.Ratio	-1.695197e-01
perc.alumni	2.961645e-01	perc.alumni	1.907969e-01
Expend	-3.614924e-04	Expend	-2.280042e-06

<code>> coef(ols_ridge)</code>						
(Intercept)	PrivateYes	Apps	Accept	Enroll	Top10perc	
3.989851e+01	4.853194e+00	1.737966e-03	-8.479738e-04	-9.953913e-04	-3.837352e-03	
Top25perc	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	
1.834768e-01	6.592359e-05	-1.290497e-03	7.473591e-04	2.044515e-03	-2.903097e-04	
Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend	
-2.969186e-03	1.574857e-01	-1.595563e-01	-2.197230e-01	3.290068e-01	-5.437189e-04	

It appears from the three models that there is not much difference between all the variables. The most interesting aspect of the models is that they look similar even when there is no regularization applied.

5. & 6. The ridge regression fit model is then looked at against the training and test sets to see how well it performs. This is shown by calculating the root mean square error for each pair as well as the root mean square error of the full model.

```

> #Compare rmse values
> ridge_rmse_train
[1] 12.96324
> ridge_rmse_test
[1] 12.9456
> rmse(test$Grad.Rate, preds_ols_ridge)
[1] 13.30244

```

From these fit model values, the training and test RMSE values are quite close. This suggests that the model is performing similarly on both the training and test data, which is a good sign. Compared to the full model the values are also close, so it does not look like overfitting from the training and test sets is a concern.


Analysis: Lasso Regression

7. & 8. The lasso regression model will be built first by using the `cv.glmnet` function to establish the `lambda.min` and `lambda.1se` values and then will be put into a plot.

```

> set.seed(123)
> cv.lasso <- cv.glmnet(train_x, train_y, alpha = 1, nfolds = 10)
> log(cv.lasso$lambda.min)
[1] -2.14716
> log(cv.lasso$lambda.1se)
[1] -0.007383546

```



The lambda.min and lambda.1se values are both small, and when plotted the lambda.min value has removed three variables from the table and the lambda.1se value has removed eight variables.

9. Implementing the cv.glmnet function, a lasso regression model is fitted, and the coefficients are compared between lambda.min and lambda.1se followed by the coefficients with no regularization.

<code>> coef(model_lasso)</code>		<code>> coef(model_1se_lasso)</code>	
18 x 1 sparse Matrix of class "dgCMatrix"		18 x 1 sparse Matrix of class "dgCMatrix"	
	s0		s0
(Intercept)	38.1215088695	(Intercept)	36.4479897891
PrivateYes	4.6092706321	PrivateYes	1.8353316358
Apps	0.0010515998	Apps	0.0002700320
Accept	.	Accept	.
Enroll	-0.0004925081	Enroll	.
Top10perc	0.0159020922	Top10perc	0.0098867060
Top25perc	0.1749512633	Top25perc	0.1678576382
F.Undergrad	.	F.Undergrad	.
P.Undergrad	-0.0012492900	P.Undergrad	-0.0007608107
Outstate	0.0006818643	Outstate	0.0007425247
Room.Board	0.0020495405	Room.Board	0.0016429893
Books	.	Books	.
Personal	-0.0029337997	Personal	-0.0021294365
PhD	0.1059159235	PhD	.
Terminal	-0.1090245955	Terminal	.
S.F.Ratio	-0.1559738443	S.F.Ratio	.
perc.alumni	0.3194285107	perc.alumni	0.2686584846
Expend	-0.0003837549	Expend	.

```
> coef(ols_lasso)
(Intercept) PrivateYes Apps Accept Enroll Top10perc
3.989851e+01 4.853194e+00 1.737966e-03 -8.479738e-04 -9.953913e-04 -3.837352e-03
Top25perc F.Undergrad P.Undergrad Outstate Room.Board Books
1.834768e-01 6.592359e-05 -1.290497e-03 7.473591e-04 2.044515e-03 -2.903097e-04
Personal PhD Terminal S.F.Ratio perc.alumni Expend
-2.969186e-03 1.574857e-01 -1.595563e-01 -2.197230e-01 3.290068e-01 -5.437189e-04
```

Between the lambda models there have been some variables that have been eliminated, with the variables 'Accept', 'F. Underground', and 'Books' being removed from both models. Overall, the coefficients from the lambda.min and lambda.1se models are both a lot smaller than the fitted ridge regression models with a majority of the coefficients being very close to zero.

10. & 11. The lasso regression fit model is then evaluated against the training and test sets to see how well they perform. This is done by calculating the root mean square error for each pair as well as the root mean square error of the full model.

```

> #Compare rmse values
> train_rmse_lasso
[1] 12.87856
> test_rmse_lasso
[1] 12.91357
> rmse(test$Grad.Rate, preds_ols_lasso)
[1] 13.30244
>

```

Similar to the root mean square error values of the ridge regression fit models, the training and test RMSE values are also quite close. This suggests that the model is performing similarly on both the training and test data. Additionally, both training and test set RMSE values are lower than the RMSE of the full model. These observations indicate that the model is not exhibiting significant overfitting.

12. Comparing the ridge and lasso models to the test set, their root mean square error values are alike. They would both perform well for this data set, but if one regression model needed to be chosen then the lasso RMSE (12.91357) would get the edge over the ridge RMSE (12.9456).

13. The stepwise selection is performed to compare against the ridge and lasso regression models followed by a fitted model to determine whether this method is an improvement.

```

> summary(model_step)

Call:
lm(formula = Grad.Rate ~ Private + Apps + Top25perc + P.Undergrad +
    Outstate + Room.Board + Personal + PhD + Terminal + perc.alumni +
    Expend, data = College)

Residuals:
    Min       1Q   Median       3Q      Max
-51.684  -7.488  -0.282   7.363  53.482

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  33.4888648   3.3489573   10.000  < 2e-16 ***
PrivateYes    3.5847682   1.6283712    2.201  0.02800 *
Apps          0.0008950   0.0001609    5.563  3.67e-08 ***
Top25perc     0.1697318   0.0321993    5.271  1.76e-07 ***
P.Undergrad  -0.0016749   0.0003631   -4.613  4.65e-06 ***
Outstate      0.0010061   0.0002257    4.458  9.51e-06 ***
Room.Board    0.0018799   0.0005795    3.244  0.00123 **
Personal     -0.0018516   0.0007485   -2.474  0.01358 *
PhD           0.0997365   0.0554704    1.798  0.07257 .
Terminal     -0.0950484   0.0612000   -1.553  0.12082
perc.alumni   0.2887259   0.0484841    5.955  3.96e-09 ***
Expend       -0.0003942   0.0001290   -3.055  0.00233 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 12.73 on 765 degrees of freedom
Multiple R-squared:  0.4585,    Adjusted R-squared:  0.4507
F-statistic: 58.88 on 11 and 765 DF,  p-value: < 2.2e-16

```

```

> coef(model_step)
      (Intercept) PrivateYes      Apps      Top25perc  P.Undergrad      Outstate
33.4888648432    3.5847682354  0.0008949751  0.1697317832 -0.0016748544  0.0010060810
      Room.Board      Personal      PhD      Terminal  perc.alumni      Expend
0.0018799059 -0.0018516261  0.0997365432 -0.0950484084  0.2887259430 -0.0003942095
> stepwise_rmse
[1] 12.74155

```

The calculated root mean square error of 12.74155 is a small improvement over the lasso regression model. The coefficients of the variables that remain after the model is performed are also closer to zero, and the intercept of 33.48 from the table is smaller than the lasso and ridge regression models. I would prefer to use this regression model over the previous ones because of these factors along with the benefits of stepwise selection that include reducing the risk of overfitting and its ability to improve the model's overall performance.

Conclusion

The assignment covered several regularization methods on a dataset to see which regression model would perform the best on a selected dependent variable. Each function used had reliable results after the training and test sets from the data set were implemented and all showed evidence of the data having a lower chance of overfitting. My choice of the regression models would be the stepwise selection for having the lowest root mean square error and for having the capability to identify the right predictors for the response variable.

Appendix

```
1 #Module 4 Assignment - Regularization
2
3 install.packages("ISLR")
4 library(ISLR)
5 install.packages("pROC")
6 library(pROC)
7 install.packages("caret")
8 library(caret)
9 library(ggplot2)
10 library(gridExtra)
11 install.packages("glmnet")
12 library(glmnet)
13 install.packages("Metrics")
14 library(Metrics)
15 attach(College)
16
17 #Question 1: Split the data into a train and test set - refer to the Feature_Selection_R.pdf
18 #document for information on how to split a data set.
19
20 #Load data
21 data(College)
22 dataset <- College
23
24 #Split data into train and test sets
25 set.seed(123)
26 trainIndex <- sort(sample(x = nrow(dataset), size = nrow(dataset) * 0.7))
27 train <- dataset[trainIndex,]
28 test <- dataset[-trainIndex,]
29
30 train_x <- model.matrix(Grad.Rate ~., train)[,-1]
31
32 test_x <- model.matrix(Grad.Rate ~., test)[,-1]
33
34 train_y <- train$Grad.Rate
35 test_y <- test$Grad.Rate
36
37 head(train_x, n = 10)
38 head(test_x, n = 10)
39 head(train_y, n = 10)
40 head(test_y, n = 10)
41
42 #Question 2: Use the cv.glmnet function to estimate the lambda.min and lambda.1se values. Compare and
43 #discuss the values.
44 set.seed(123)
45 cv.ridge <- cv.glmnet(train_x, train_y, alpha = 0, nfolds = 10)
46
47 log(cv.ridge$lambda.min)
48 log(cv.ridge$lambda.1se)
49
50 #Question 3: Plot the results from the cv.glmnet function provide an interpretation. What does this plot
51 #tell us?
52 plot(cv.ridge)
53
54 #Question 4: Fit a Ridge regression model against the training set and report on the coefficients. Is there
55 #anything interesting?
56 model.ridge <- glmnet(train_x, train_y, alpha = 0, lambda = cv.ridge$lambda.min)
57
58 #Display regression coefficients
59 coef(model.ridge)
60
```



```

61 #Fit the final model on the training data using lambda.lse
62 model_lse_ride <- glmnet(train_x, train_y, alpha = 0, lambda = cv.ride$lambda.lse)
63
64 #Display regression coefficients
65 coef(model_lse_ride)
66
67 #Display coefficients of ols model with no regularization
68 ols_ride <- lm(Grad.Rate ~ ., data = train)
69 coef(ols_ride)
70
71 #Question 5: Determine the performance of the fit model against the training set by calculating the root
72 #mean square error (RMSE). sqrt(mean((actual - predicted)^2))
73
74 #View RMSE of full model
75 preds_ols_ride <- predict(ols_ride, new = test)
76 rmse(test$Grad.Rate, preds_ols_ride)
77
78 #Make predictions on the train data using lambda.min
79 preds_train_ride <- predict(model_lse_ride, newx = train_x)
80 train_rmse_ride <- rmse(train_y, preds_train_ride)
81
82 #Compute the RMSE
83 ridge_rmse_train <- sqrt(mean((train_y - preds_train_ride)^2))
84 ridge_rmse_train
85
86 #Question 6: Determine the performance of the fit model against the test set by calculating the root mean
87 #square error (RMSE). Is your model overfit?
88
89 preds_test_ride <- predict(model_lse_ride, newx = test_x)
90 test_rmse_ride <- rmse(test_y, preds_test_ride)

```

```

91
92 #Compute the RMSE
93 ridge_rmse_test <- sqrt(mean((test_y - preds_test_ride)^2))
94 ridge_rmse_test
95
96 #Compare rmse values
97 ridge_rmse_train
98 ridge_rmse_test
99 rmse(test$Grad.Rate, preds_ols_ride)
100
101 #Question 7: Use the cv.glmnet function to estimate the lambda.min and lambda.lse values. Compare and
102 #discuss the values.
103 set.seed(123)
104 cv.lasso <- cv.glmnet(train_x, train_y, alpha = 1, nfolds = 10)
105
106 log(cv.lasso$lambda.min)
107 log(cv.lasso$lambda.lse)
108
109 #Question 8: Plot the results from the cv.glmnet function provide an interpretation. What does this plot
110 #tell us?
111 plot(cv.lasso)
112
113 #Question 9: Fit a LASSO regression model against the training set and report on the coefficients. Do any
114 #coefficients reduce to zero? If so, which ones?
115 model_lasso <- glmnet(train_x, train_y, alpha = 1, lambda = cv.lasso$lambda.min)
116 model_lasso
117
118 #Display regression coefficients
119 coef(model_lasso)
120

```

```

121 #Fit the final model on the training data using lambda.lse
122 model_lse_lasso <- glmnet(train_x, train_y, alpha = 1, lambda = cv.lasso$lambda.lse)
123
124 #Display regression coefficients
125 coef(model_lse_lasso)
126
127 #Display coefficients of ols model with no regularization
128 ols_lasso <- lm(Grad.Rate ~ ., data = train)
129 coef(ols_lasso)
130
131 #Question 10: Determine the performance of the fit model against the training set by calculating the root
132 #mean square error (RMSE). sqrt(mean((actual - predicted)^2))
133
134 #View RMSE of full model
135 preds_ols_lasso <- predict(ols_lasso, new = test)
136 rmse(test$Grad.Rate, preds_ols_lasso)
137
138 #Make predictions on the train data using lambda.min
139 preds_train_lasso <- predict(model_lse_lasso, newx = train_x)
140 train_rmse_lasso <- rmse(train_y, preds_train_lasso)
141
142 #Compute the RMSE
143 lasso_rmse <- sqrt(mean((train_y - preds_train_lasso)^2))
144 lasso_rmse
145
146 #Question 11: Determine the performance of the fit model against the test set by calculating the root mean
147 #square error (RMSE). Is your model overfit?
148
149 preds_test_lasso <- predict(model_lse_lasso, newx = test_x)
150 test_rmse_lasso <- rmse(test_y, preds_test_lasso)

```

```

150 test_rmse_lasso <- rmse(test_y, preds_test_lasso)
151
152 #Compare rmse values
153 train_rmse_lasso
154 test_rmse_lasso
155 rmse(test$Grad.Rate, preds_ols_lasso)
156
157 #Comparison
158 #Question 12: Which model performed better and why? Is that what you expected?
159
160 #Question 13: Refer to the Intermediate_Analytics_Feature_Selection_R.pdf document for how to perform stepwise
161 #selection and then fit a model. Did this model perform better or as well as Ridge regression
162 #or LASSO? Which method do you prefer and why?
163
164 # Stepwise selection method and fit the model
165 step<lm(Grad.Rate ~ ., data = College), direction = 'both')
166 model_step <- step<lm(Grad.Rate ~ ., data = College), direction = 'both')
167 summary(model_step)
168 coef(model_step)
169
170 #Make predictions on Grad Rate
171 stepwise <- predict(model_step, newdata = test)
172 stepwise
173
174 #Compute the RMSE
175 stepwise_rmse <- sqrt(mean((test$Grad.Rate - stepwise)^2))
176 stepwise_rmse
177

```

References

Barkved, K. (2022, Feb. 11th). *The Difference Between Training Data vs. Test Data in Machine Learning*. Obviously.ai. <https://www.obviously.ai/post/the-difference-between-training-data-vs-test-data-in-machine-learning>

Breheny, P. & Zeng, Y. (2018). The biglasso Package: A Memory- and Computation-Efficient Solver for Lasso Model Fitting with Big Data in R. *Journal of Statistical Software, Volume VV, Issue II*. <https://arxiv.org/pdf/1701.05936.pdf>

Chat GPT. (2023, December 10th). Default (GPT 3.5). < <https://chat.openai.com/>>